

Domain specific computing architectures - II

Asif Ali Khan

Fall Semester 2024

Department of Computer Systems Engineering


UET Peshawar, Pakistan

Oct 24, 2024

Recap: Programmability of Von Neumann systems


- ❑ ISA is the set of instruction the processor can understand and is the interface between h/w and s/w
- ❑ High level languages are typically used to program these systems
- ❑ Compilers transform this hardware-agnostic representation to machine code
- ❑ A compiler typically has front-end, middle-end, and back-end

Recap: why are DSAs needed?



Aguirre et al, Nature Communications, 2024

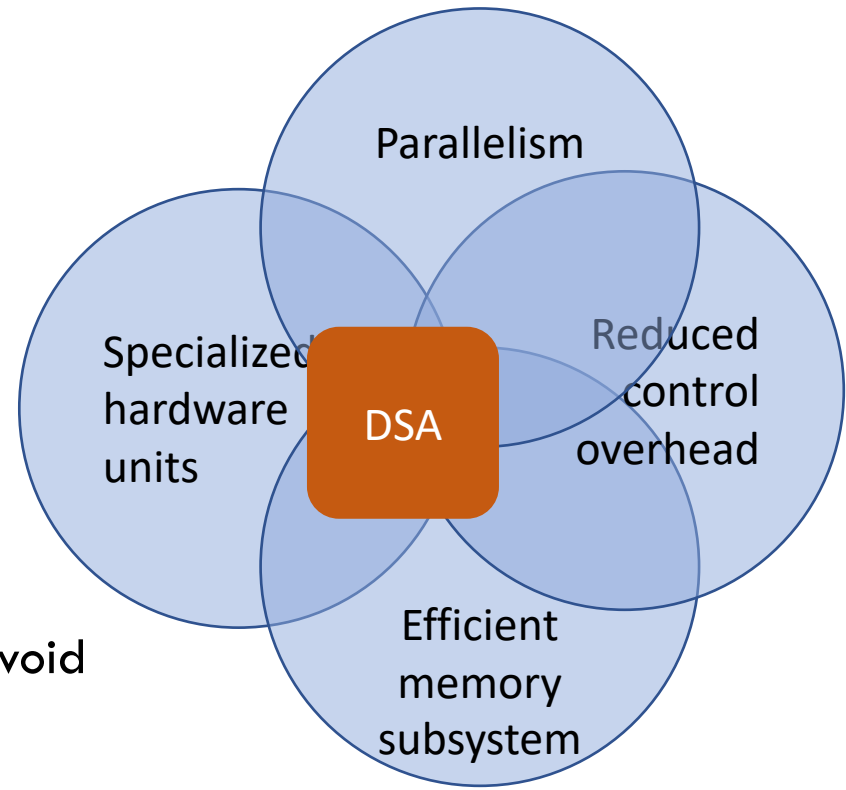
Traditional general-purpose (GP) computing can not fulfill the compute-demands of these applications due to their inefficiencies



Dally et al, Efficient embedded computing, IEEE'08

DSAs: Important aspects

- ❑ Specialization can be in:
 - ❑ Hardware
 - ❑ Data
- ❑ Parallelism can be decided as per the domain requirements
- ❑ Efficient local memories can be utilized to avoid long data access latencies
- ❑ The control overhead of over 90% in CPUs can be significantly reduced



DSAs: Specialization

- ❑ Specialization can be in:
 - ❑ Data
 - ❑ Hardware

DSAs: Specialization

- ❑ Specialization can be in:
 - ❑ Data
 - ❑ Hardware
- ❑ Data specialization: Instead of using the general-purpose 32/64 bits FP, use specialized types for different domains (for ML, 1-8 bits)

DSAs: Specialization

- ❑ Specialization can be in:
 - ❑ Data
 - ❑ Hardware
- ❑ Data specialization: Instead of using the general-purpose 32/64 bits FP, use specialized types for different domains (for ML, 1-8 bits)
- ❑ Hardware specialization: Specialized hardware modules, e.g., multiply-and-accumulate (MAC) unit, adder tree, other special function units (e.g., for softmax)

DSAs: Specialization

$$I(i, j) = \max \begin{cases} H(i, j-1) - o \\ I(i, j-1) - e \end{cases}$$

$$D(i, j) = \max \begin{cases} H(i-1, j) - o \\ D(i-1, j) - e \end{cases}$$

$$H(i, j) = \max \begin{cases} 0 \\ I(i, j) \\ D(i, j) \\ H(i-1, j-1) + W(r_i, q_j) \end{cases}$$

- Smith-Waterman algorithm for sequence alignment

Reference

		*	G	C	G	A	C	T	T	T
	*	0	0	0	0	0	0	0	0	0
	G	0	2	-1	2	-1	0	0	0	0
	T	0	1	1	1	1	0	2	2	2
	C	0	0	3	-2	-1	3	-2	1	1
	G	0	2	2	5	4	3	2	1	0
	T	0	1	1	4	4	3	5	4	3
	T	0	0	0	3	3	3	5	7	6
	T	0	0	0	2	2	2	5	7	9

Query

DSAs: Specialization

$$I(i, j) = \max \begin{cases} H(i, j-1) - o \\ I(i, j-1) - e \end{cases}$$

$$D(i, j) = \max \begin{cases} H(i-1, j) - o \\ D(i-1, j) - e \end{cases}$$

$$H(i, j) = \max \begin{cases} 0 \\ I(i, j) \\ D(i, j) \\ H(i-1, j-1) + W(r_i, q_j) \end{cases}$$

- ❑ Smith-Waterman algorithm for sequence alignment
- ❑ 15 loads/stores, 35 arithmetic/logic ops

DSAs: Specialization

$$I(i, j) = \max \begin{cases} H(i, j-1) - o \\ I(i, j-1) - e \end{cases}$$

$$D(i, j) = \max \begin{cases} H(i-1, j) - o \\ D(i-1, j) - e \end{cases}$$

$$H(i, j) = \max \begin{cases} 0 \\ I(i, j) \\ D(i, j) \\ H(i-1, j-1) + W(r_i, q_j) \end{cases}$$

- ❑ Smith-Waterman algorithm for sequence alignment
- ❑ 15 loads/stores, 35 arithmetic/logic ops
- ❑ Intel Xeon 14nm:
 - ❑ 37 cycles (latency), 81nj (energy)

DSAs: Specialization

$$I(i, j) = \max \begin{cases} H(i, j-1) - o \\ I(i, j-1) - e \end{cases}$$


$$D(i, j) = \max \begin{cases} H(i-1, j) - o \\ D(i-1, j) - e \end{cases}$$

$$H(i, j) = \max \begin{cases} 0 \\ I(i, j) \\ D(i, j) \\ H(i-1, j-1) + W(r_i, q_j) \end{cases}$$

Turakhia, Y., Bejerano, G. and Dally, W.J., Darwin: A Genomics Co-processor Provides up to 15,000 X Acceleration on Long Read Assembly. ASPLOS, 2018

- ❑ Smith-Waterman algorithm for sequence alignment
- ❑ 15 loads/stores, 35 arithmetic/logic ops
- ❑ Intel Xeon 14nm:
 - ❑ 37 cycles (latency), 81nj (energy)
- ❑ Darwin accelerator:
 - ❑ Latency: 1 cycle (37x speedup)
 - ❑ Energy: 3.1pj (26,000x reduction)

DSAs: Parallelism



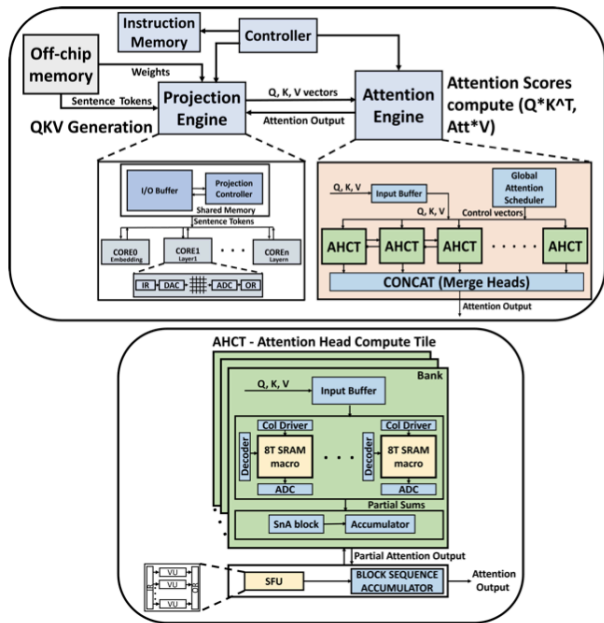
- ❑ One of the fundamental principals of DSAs

DSAs: Parallelism

- ❑ One of the fundamental principals of DSAs
- ❑ Parallelism is usually hierarchal, e.g., multiple MAC units per processing element (PE), multiple PEs per module, multiple modules per system

DSAs: Parallelism

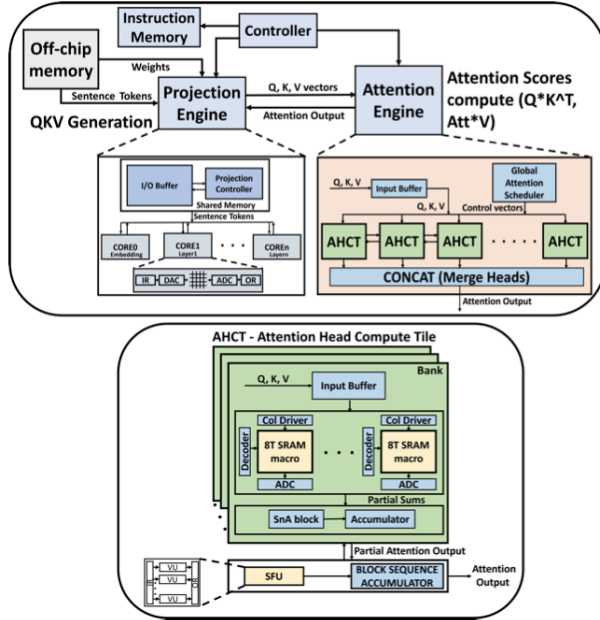
- ❑ One of the fundamental principals of DSAs
- ❑ Parallelism is usually hierarchal, e.g., multiple MAC units per processing element (PE), multiple PEs per module, multiple modules per system



S. Sridharan et al., "X-Former: In-Memory Acceleration of Transformers", IEEE TVLSI 2023

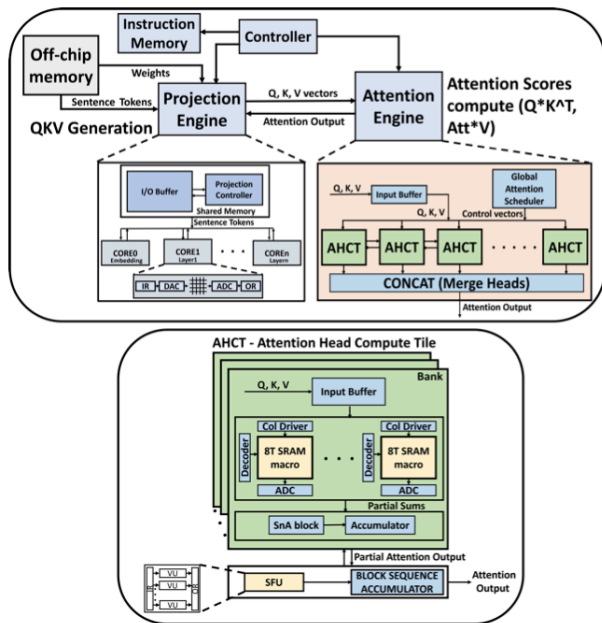
DSAs: Parallelism

- ❑ One of the fundamental principals of DSAs
- ❑ Parallelism is usually hierarchal, e.g., multiple MAC units per processing element (PE), multiple PEs per module, multiple modules per system
- ❑ Memory organization is key to performance



S. Sridharan et al., "X-Former: In-Memory Acceleration of Transformers", IEEE TVLSI 2023


DSAs: Parallelism



- ❑ One of the fundamental principals of DSAs
- ❑ Parallelism is usually hierarchal, e.g., multiple MAC units per processing element (PE), multiple PEs per module, multiple modules per system
- ❑ Memory organization is key to performance
- ❑ Parallel PEs must exploit locality

S. Sridharan et al., "X-Former: In-Memory Acceleration of Transformers", IEEE TVLSI 2023

DSAs: Parallelism




- ❑ One of the fundamental principals of DSAs
- ❑ Parallelism is usually hierarchal, e.g., multiple MAC units per processing element (PE), multiple PEs per module, multiple modules per system
- ❑ Memory organization is key to performance
- ❑ Parallel PEs must exploit locality
- ❑ There should be ideally no cross-PE dependencies (we will see this in the UPMEM example)

S. Sridharan et al., "X-Former: In-Memory Acceleration of Transformers", IEEE TVLSI 2023

DSAs: Parallelism in Darwin


- ❑ Outer-loop: 64 PEs running alignment instances in parallel
 - ❑ No cross-PEs communication

DSAs: Parallelism in Darwin




- ❑ Outer-loop: 64 PEs running alignment instances in parallel
 - ❑ No cross-PEs communication
- ❑ Inner-loop: In each one of them, 64 special function units compute H, I, Ds in parallel
 - ❑ Only neighbors communicate

DSAs: Parallelism in Darwin



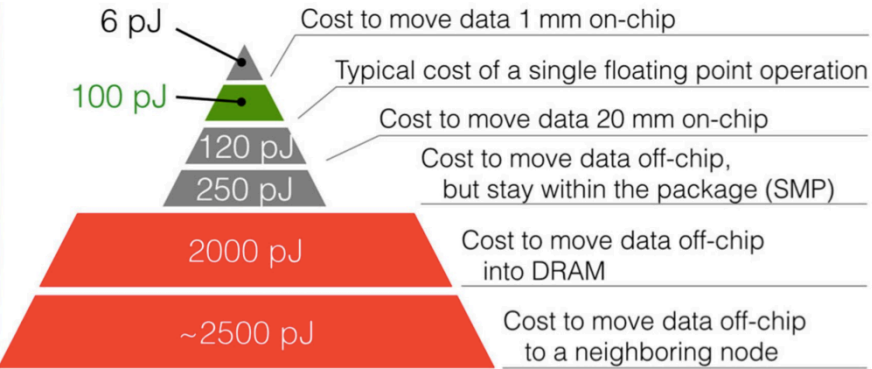
- ❑ Outer-loop: 64 PEs running alignment instances in parallel
 - ❑ No cross-PEs communication
- ❑ Inner-loop: In each one of them, 64 special function units compute H, I, Ds in parallel
 - ❑ Only neighbors communicate
- ❑ Speedup:
 - ❑ Specialization: 37x
 - ❑ Parallelization: 4034x
 - ❑ Total speedup: 150,000x

DSAs: Optimized memory



- Recall from the last lecture:
Arithmetic is free; memory is expensive

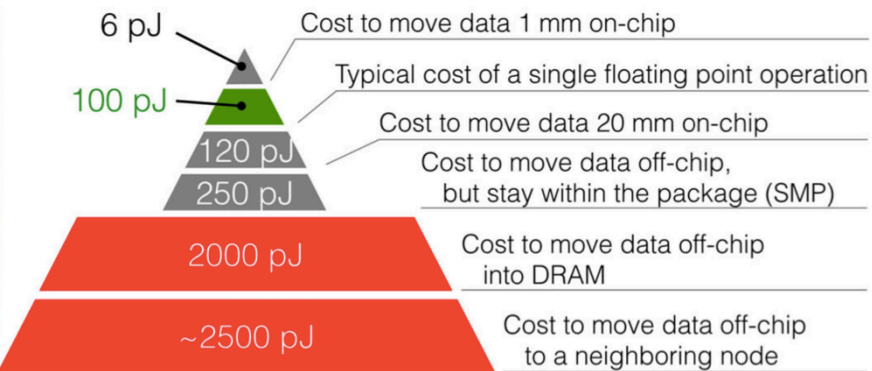
DSAs: Optimized memory



You, Y., et al. Fast LSTM by dynamic decomposition on cloud and distributed systems. Knowledge and Information Systems, 2020

- ❑ Recall from the last lecture that arithmetic is almost free; memory is expensive
- ❑ Data movement is prohibitively expensive


DSAs: Optimized memory



You, Y., et al. Fast LSTM by dynamic decomposition on cloud and distributed systems. Knowledge and Information Systems, 2020

- ❑ Recall from the last lecture that arithmetic is almost free; memory is expensive
- ❑ Data movement is prohibitively expensive
- ❑ Key: Use many small, local memory to
 - ❑ Minimize distance
 - ❑ Maximize reuse


DSAs: Optimized memory



You, Y., et al. Fast LSTM by dynamic decomposition on cloud and distributed systems. Knowledge and Information Systems, 2020


- ❑ Recall from the last lecture that arithmetic is almost free; memory is expensive
- ❑ Data movement is prohibitively expensive
- ❑ Key: Use many small, local memory to
 - ❑ Minimize distance
 - ❑ Maximize reuse
- ❑ Other optimizations:
 - ❑ Data compression
 - Increase the effective bandwidth/ capacity

DSAs: Reduced overhead




- ❑ Specialized hardware reduces overhead

DSAs: Reduced overhead




- ❑ Specialized hardware reduces overhead
- ❑ In-order general purpose CPU spends >90% energy on overhead (instruction fetch, decode, data, control etc)

DSAs: Reduced overhead



- ❑ Specialized hardware reduces overhead
- ❑ In-order general purpose CPU spends >90% energy on overhead (instruction fetch, decode, data, control etc)
- ❑ In OOO processor, this is over 99%

DSAs: Reduced overhead



- ❑ Specialized hardware reduces overhead
- ❑ In-order general purpose CPU spends >90% energy on overhead (instruction fetch, decode, data, control etc)
- ❑ In OOO processor, this is over 99%
- ❑ Example:
 - ❑ ARM A-15, integer add: 250pj
 - ❑ 32-bit CMOS adder: 68fj (4000x less)

GPU architecture and programming

- ❑ Extreme throughput-oriented processors

GPU architecture and programming

- ❑ Extreme throughput-oriented processors
- ❑ The basic computing unit of a GPU is called a CUDA core

GPU architecture and programming

- ❑ Extreme throughput-oriented processors
- ❑ The basic computing unit of a GPU is called a CUDA core
- ❑ CUDA cores are combined to form streaming multi-processors (SMs)

GPU architecture and programming

- ❑ Extreme throughput-oriented processors
- ❑ The basic computing unit of a GPU is called a CUDA core
- ❑ CUDA cores are combined to form streaming multi-processors (SMs)
- ❑ SMs are organized into grids

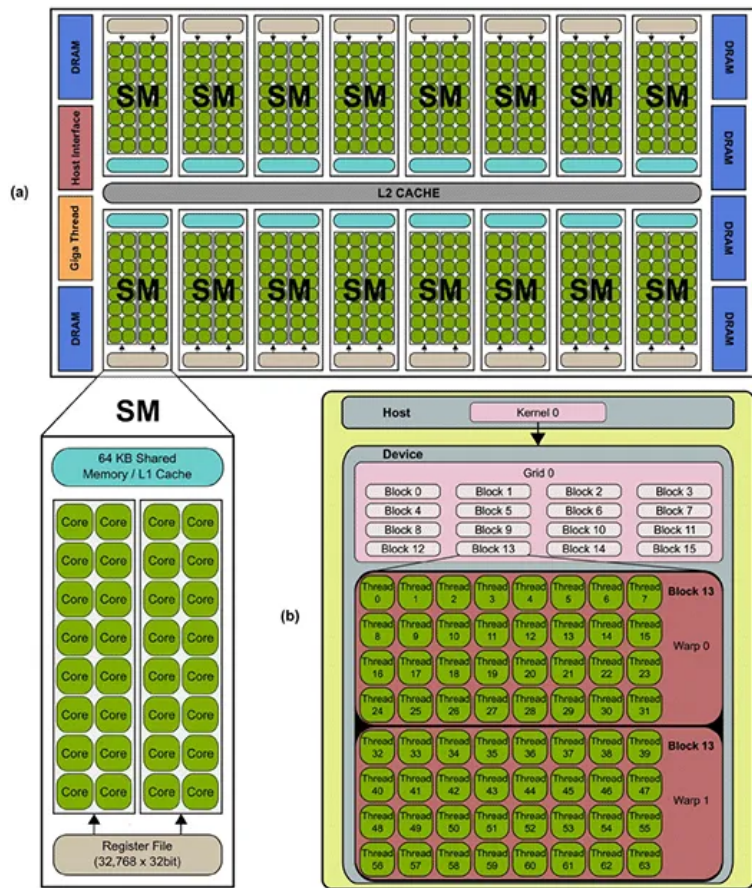
GPU architecture and programming

- ❑ Extreme throughput-oriented processors
- ❑ The basic computing unit of a GPU is called a CUDA core
- ❑ CUDA cores are combined to form streaming multi-processors (SMs)
- ❑ SMs are organized into grids
- ❑ Cores within an SM communicate via local (shared) memory

GPU architecture and programming

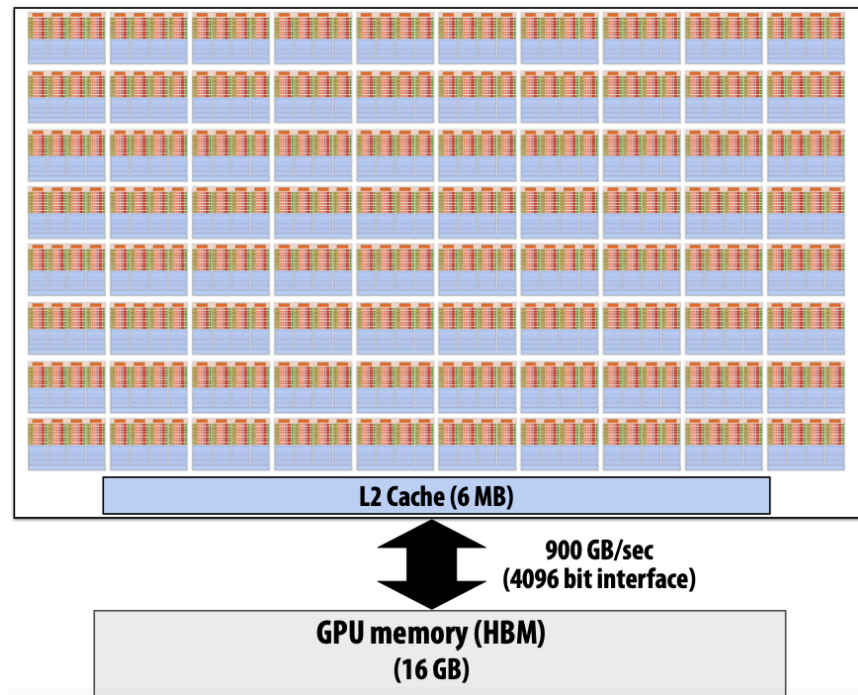
- ❑ Extreme throughput-oriented processors
- ❑ The basic computing unit of a GPU is called a CUDA core
- ❑ CUDA cores are combined to form streaming multi-processors (SMs)
- ❑ SMs are organized into grids
- ❑ Cores within an SM communicate via local (shared) memory
- ❑ SMs communicate via global (device) memory

GPU architecture and programming



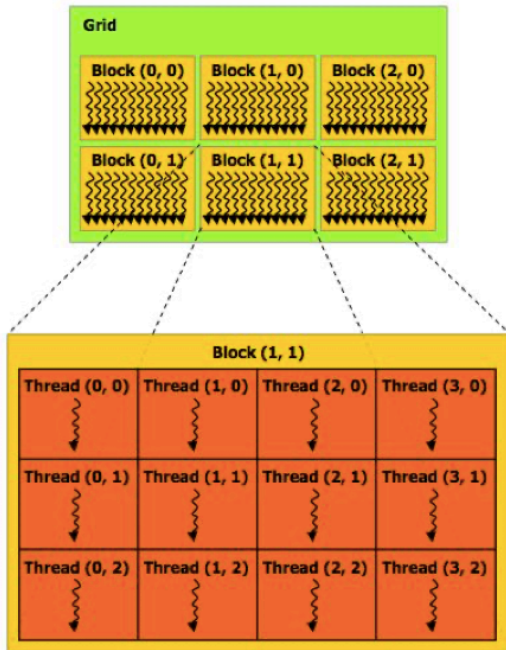
GPU architecture and programming

- ❑ Nvidia V100 GPU
- ❑ 80 SM cores
- ❑ 64 FP32 ALUs per SM core
- ❑ 5120 FP32 ALUs per board



GPU architecture and programming

- ❑ CUDA programming language is used to program Nvidia GPUs



Regular application thread running on CPU (the "host")

```
const int Nx = 12;
const int Ny = 6;

dim3 threadsPerBlock(4, 3);
dim3 numBlocks(Nx/threadsPerBlock.x, Ny/threadsPerBlock.y);

// assume A, B, C are allocated Nx x Ny float arrays

// this call will launch 72 CUDA threads:
// 6 thread blocks of 12 threads each
matrixAdd<<<numBlocks, threadsPerBlock>>>(A, B, C);
```

GPU architecture and programming

- ❑ CUDA programming language is used to program Nvidia GPUs

Regular application thread running on CPU (the “host”)

```
const int Nx = 12;
const int Ny = 6;

dim3 threadsPerBlock(4, 3);
dim3 numBlocks(Nx/threadsPerBlock.x, Ny/
threadsPerBlock.y);

// assume A, B, C are allocated Nx x Ny float arrays

// this call will launch 72 CUDA threads:
// 6 thread blocks of 12 threads each
matrixAdd<<<numBlocks, threadsPerBlock>>>(A, B, C);
```

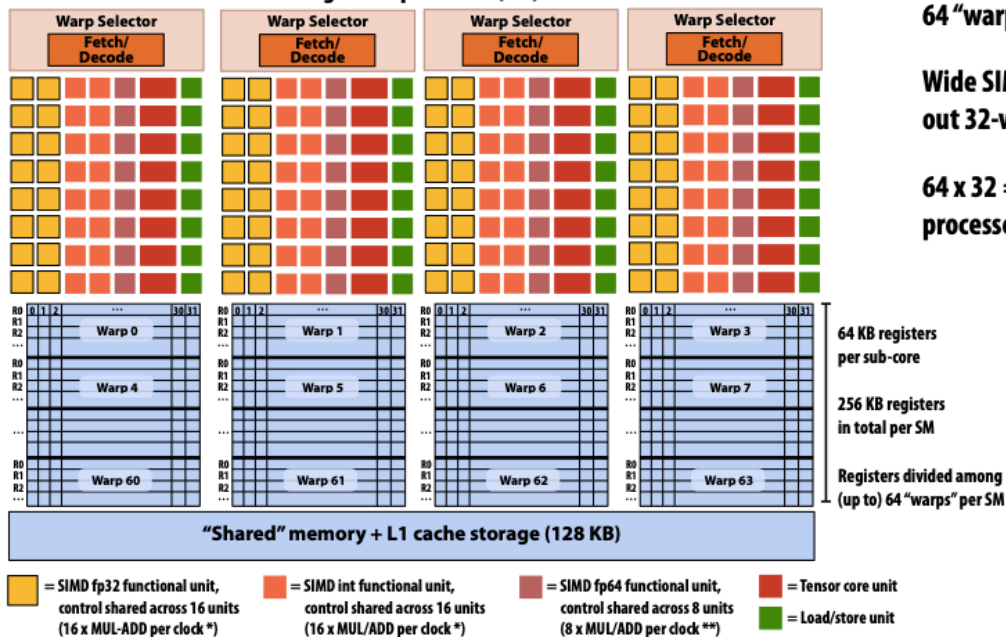
CUDA kernel definition

```
// kernel definition (runs on GPU)
__global__ void matrixAdd(float A[Ny][Nx],
                          float B[Ny][Nx],
                          float C[Ny][Nx])
{
    int i = blockIdx.x * blockDim.x + threadIdx.x;
    int j = blockIdx.y * blockDim.y + threadIdx.y;

    C[j][i] = A[j][i] + B[j][i];
}
```

GPU architecture: V100 SM unit and programming

This is one NVIDIA V100 streaming multi-processor (SM) unit



64 "warp" execution contexts per SM

Wide SIMD: 16-wide SIMD ALUs (carry out 32-wide SIMD execute over 2 clocks)

64 x 32 = up to 2048 data items processed concurrently per "SM" core

64 KB registers per sub-core

256 KB registers in total per SM

Registers divided among (up to) 64 "warps" per SM

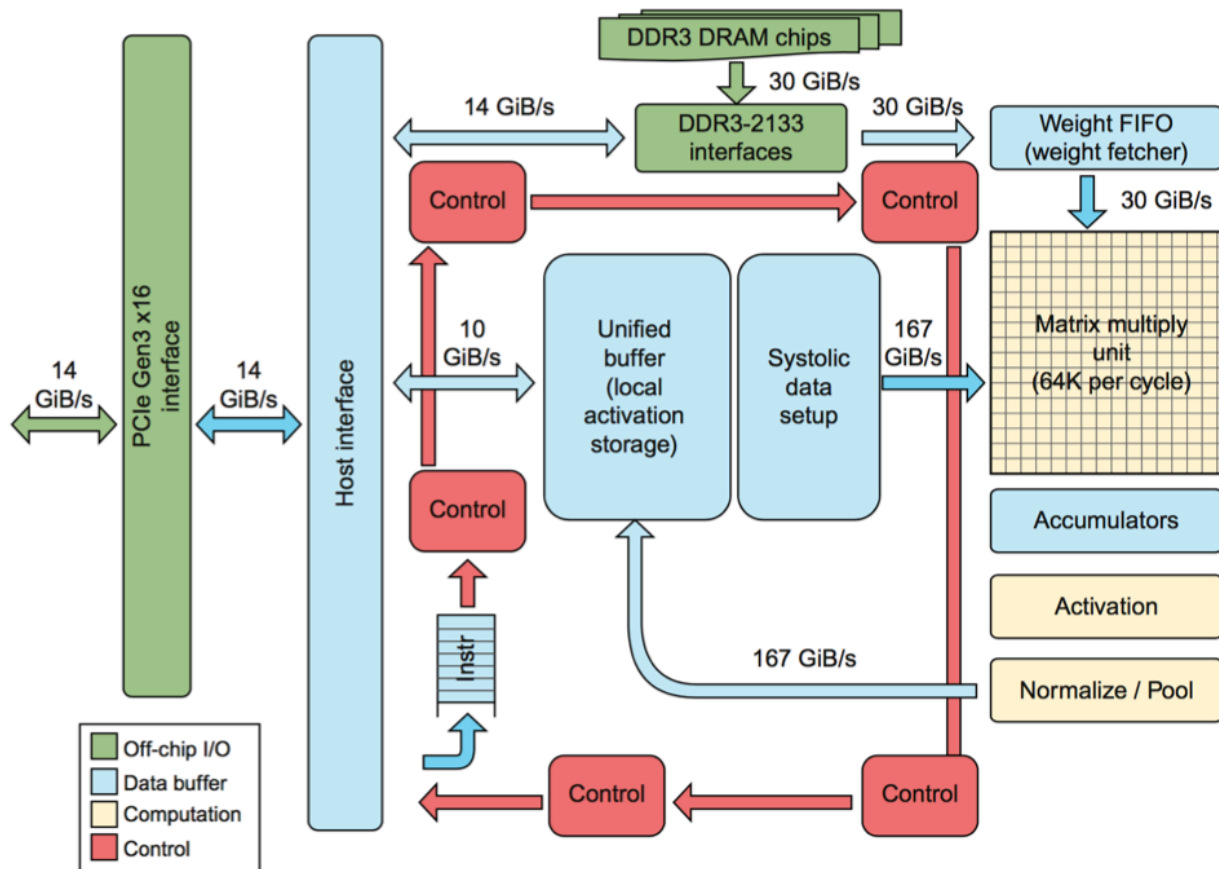
Tensor processing unit (TPU)

- ❑ Google's ASIC for deep neural networks
- ❑ Specifically designed for inference
- ❑ Is programmed using Google's DSL named TensorFlow

Key idea:

- ❑ Dedicated modules for matrix-matrix multiplication (matmul), and other functions
 - ❑ 256 x 256 MAC units
- ❑ Scratchpad memories

Tensor processing unit (TPU)



TPU's ISA

- ❑ TPU follows the CISC type, typically 10-20 cycles per instruction (CPI)

TPU's ISA

- ❑ TPU follows the CISC type, typically 10-20 cycles per instruction (CPI)
- ❑ Reduced overhead, i.e., no program counter, branch instructions etc.

TPU's ISA

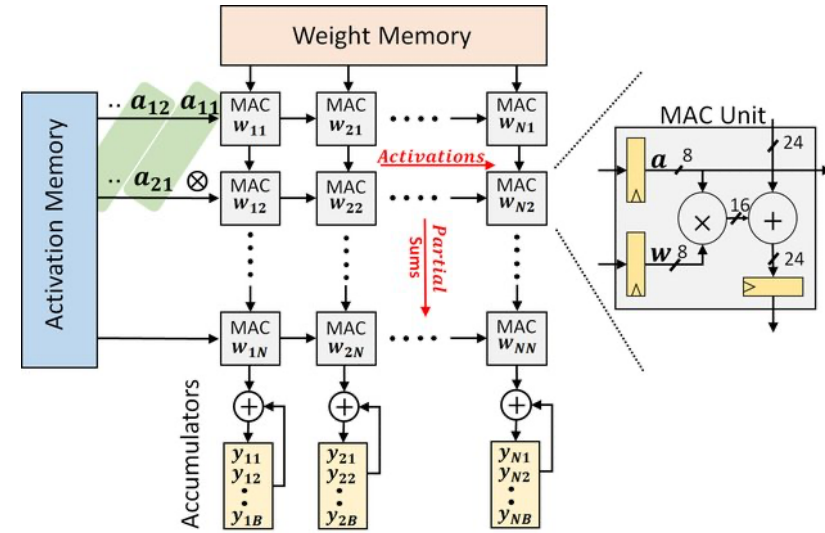
- ❑ TPU follows the CISC type, typically 10-20 cycles per instruction (CPI)
- ❑ Reduced overhead, i.e., no program counter, branch instructions etc.
- ❑ A handful (around a dozen) of instructions in total

TPU's ISA

- ❑ TPU follows the CISC type, typically 10-20 cycles per instruction (CPI)
- ❑ Reduced overhead, i.e., no program counter, branch instructions etc.
- ❑ A handful (around a dozen) of instructions in total
- ❑ Some key instructions are:
 - ❑ Read_Host_Memory: CPU memory → Unified Buffer (UB)
 - ❑ Read_Weights: Weight memory → Weight FIFO
 - ❑ MatrixMatrixMultiply/Convolve: Perform MM, MV etc; input: UB, output: accumulator
 - ❑ Activate: Performs non-linear activations, e.g., ReLU; input: accumulator, output: UB
 - ❑ Write_Host_Memory: UB → CPU memory

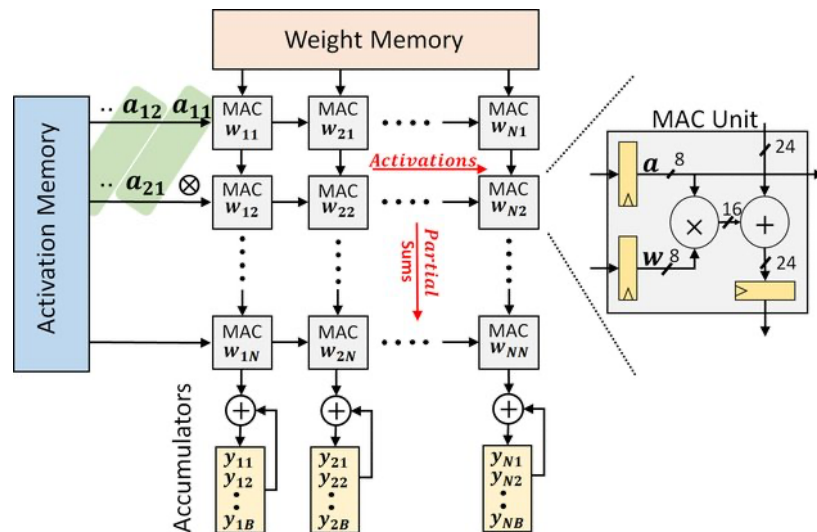
TPU's microarchitecture & summary

- TPU is based on systolic arrays
 - More details:
Kung, "Why systolic architectures?", Computers, 1982



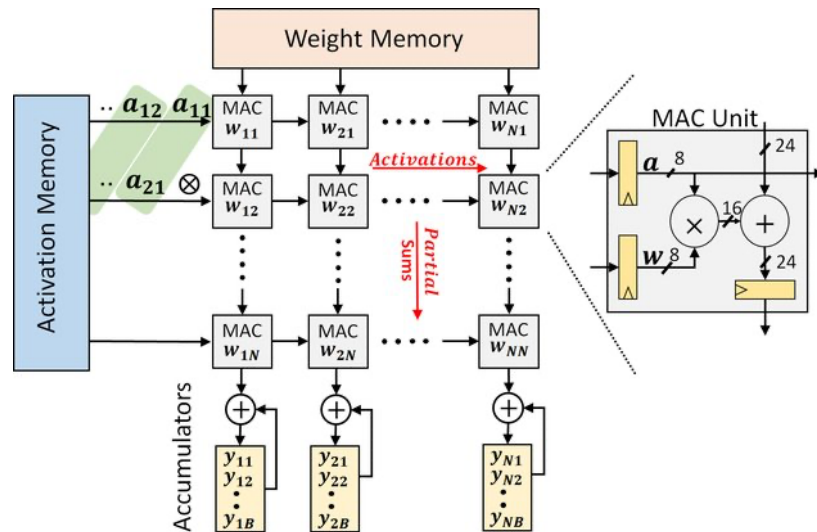
TPU's microarchitecture & summary

- TPU is based on systolic arrays
 - More details:
Kung, "Why systolic architectures?", Computers, 1982
- Uses dedicated memories
 - Dedicated buffers, accumulators




TPU's microarchitecture & summary

- TPU is based on systolic arrays
 - More details:
Kung, "Why systolic architectures?", Computers, 1982
- Uses dedicated memories
 - Dedicated buffers, accumulators
- Invest resources in arithmetic units and memories (way more resources than a server-class CPU)



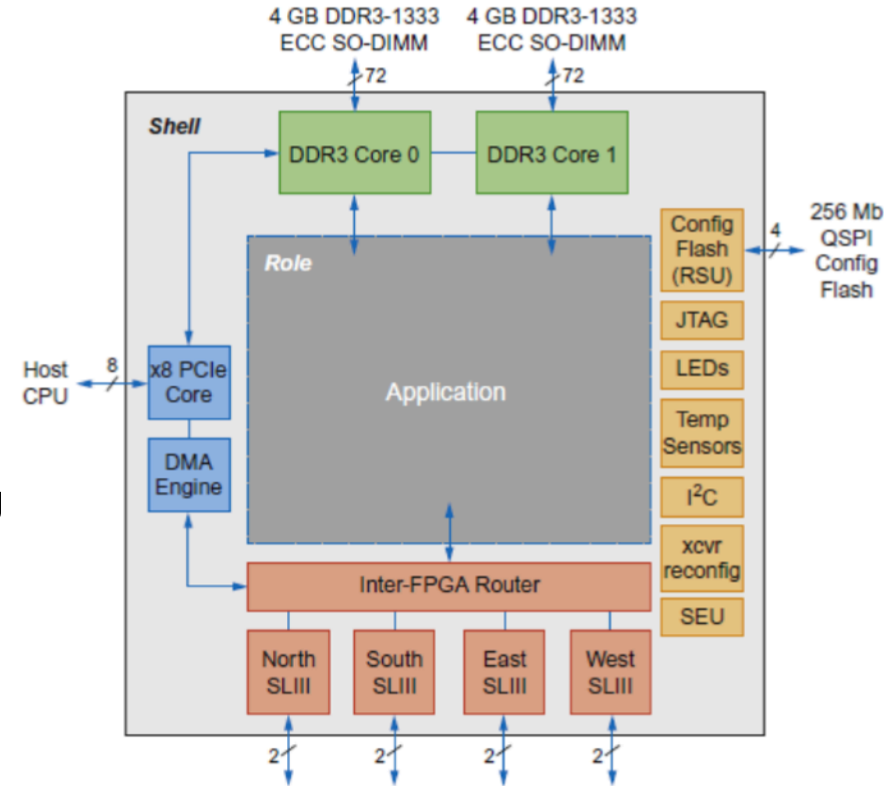
TPU's microarchitecture & summary

- ❑ TPU is based on systolic arrays
 - ❑ More details:
Kung, "Why systolic architectures?", Computers, 1982
- ❑ Uses dedicated memories
 - ❑ Dedicated buffers, accumulators
- ❑ Invest resources in arithmetic units and memories (way more resources than a server-class CPU)
- ❑ Easiest form of parallelism (SIMD, systolic arrays), and specialization (8b data)



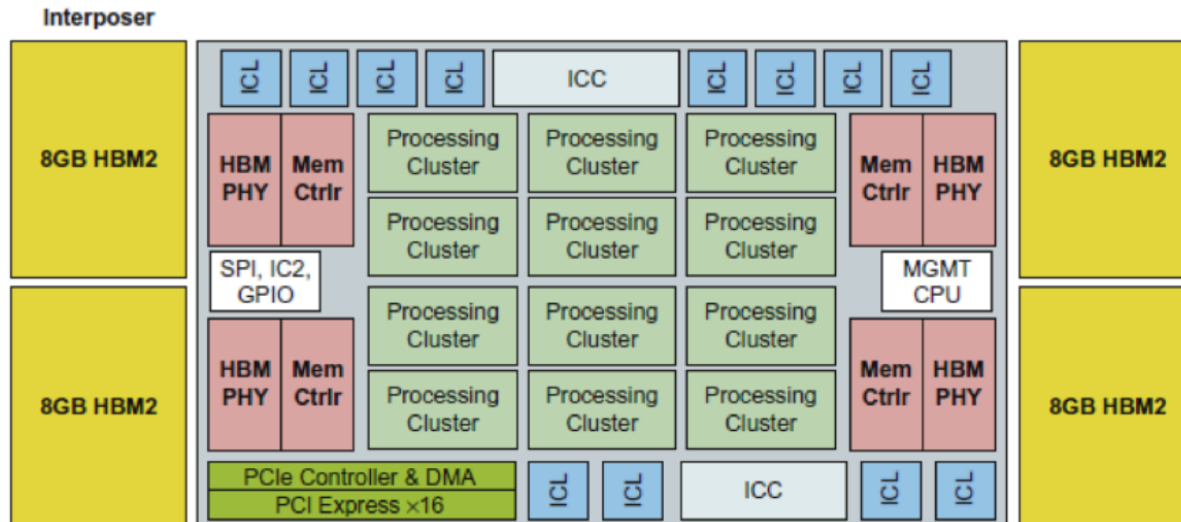
Microsoft's Catapult and Brainwave

- ❑ FPGA based DSA for AI inference (edge as well as cloud)
- ❑ Catapult: 32MB of Flash memory, PCIe connectivity
- ❑ Brainwave: Has dedicated deep-learning processing units (DPUs) soft-cores



Intel's Crest

- ❑ Specifically designed for DNNs training
- ❑ 16-bit Fixed Point
- ❑ Multiple matmul units, operating on 32 x 32 size matrices
- ❑ Employs HBM + SRAM



Apple's neural processing units (NPUs) and neural engine

- ❑ Your assignment to read on it and answer the following:
 - ❑ How is the micro-architecture?
 - ❑ How is it programmed?
 - ❑ What kind of parallelism is it using?
 - ❑ How does it compare to Google's TPU in the above aspects, and/or otherwise?

Thank you!

asif.ali@uetpeshawar.edu.pk